

Feature Extraction and Selection

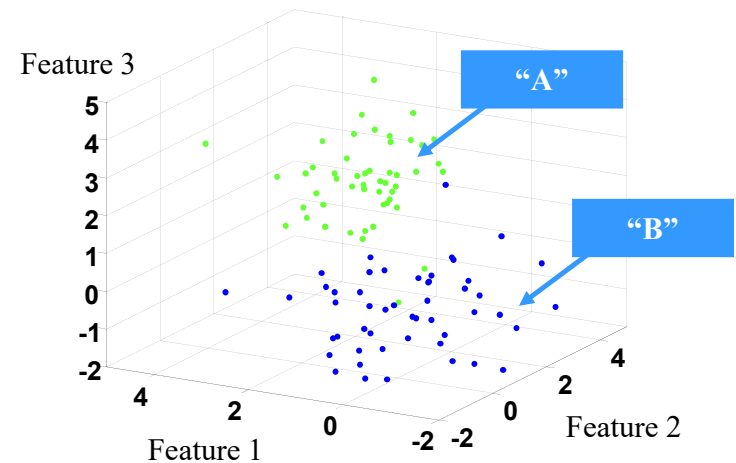
› Marco Loog

Outline

- › A bit on high-dimensional spaces
- › Some basics of feature extraction
- › Some basics of feature selection

Feature Space

- › A p -dimensional space, in which each dimension is a feature containing N [labeled] samples [objects]
- › Why and how should lower number of features?



Curse of Dimensionality


- › Problem : too few samples in too many dimensions
[the curse of dimensionality]

Let's discuss histogram-based density estimation
...with increasingly finer binning?

- › Anyway : in high-dimensional spaces,
our 2D/3D intuition does not work anymore...



High-Dimensional Spaces

High-Dimensional Spaces



- › Example : neighborhood capturing 10% of uniformly distributed data in hypercube
- › E.g. in \mathbb{R}^{20} sides of $\sqrt[20]{.1} \approx 0.89$
So, not a small block anymore...

High-Dimensional Spaces



› Example : boundary points

1000 normal samples in 2D then 1% on convex hull

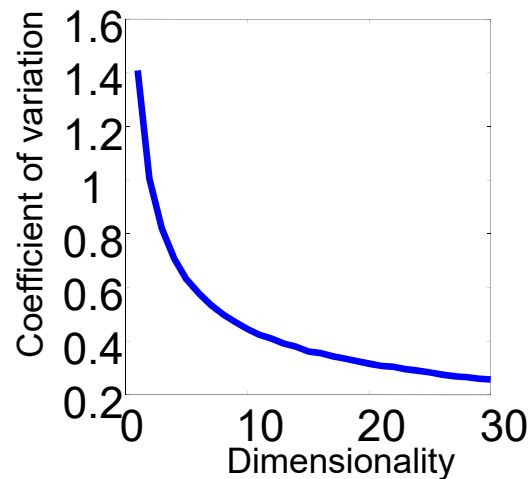
1000 in 20D then 95% on convex hull

High-Dimensional Spaces

› Example : points tend to have equal distances

Consider $\frac{\text{std}(d^2)}{\text{mean}(d^2)}$ for squared distance d^2

For points in \mathbb{R}^{1000} from standard normal, distribution is approximately $N(2000, 8000)$




High-Dimensional Spaces

- › This means [roughly] for increasing dimensionality
 - local, distance-based methods suffer most, e.g. NN-methods
 - global, more restricted models suffer less, e.g. linear models
- › So...
 - controlling classifier complexity important
 - p should be kept as low as possible : dimensionality reduction

Dimensionality Reduction by Selection and Extraction


Dimensionality Reduction

- › Problem : too few samples in too many dimensions
[the curse of dimensionality]



Dimensionality Reduction

- › Problem : too few samples in too many dimensions [the curse of dimensionality]
- › Solution : drop dimensions / features
 - Feature selection
 - Feature extraction
- › Questions :
 - Which dimensions to drop?
 - What feature subset to keep?



Dimensionality Reduction

› Other uses :


Fewer parameters give faster algorithms and parameters are easier to estimate

Explaining which measurements are useful and which are not [reducing redundancy]

Visualization of data can be a powerful tool when designing pattern recognition systems

Feature Selection vs Extraction

- › Feature selection :
select d out of
 p measurements
- › Feature extraction :
map p measurements
to d measurements



Feature Selection vs Extraction

- › Think of selection and extraction as finding a mapping
- › We need :
 - Criterion function, e.g. error, class overlap, information loss,...
 - Optimization or “search” algorithm to find mapping for given criterion

Note on Criteria

- › The optimal[?] criterion :
final performance of the entire system
Maybe calculated using cross-validation
- › Approximate performance predictors
Calculate performance of easy-to-use criterion giving indication
of how well a more powerful / realistic criterion may perform

Two Classical Linear Feature Extractors

Linear Feature Extraction

› Unsupervised :

Principal Component Analysis [PCA]

› Supervised :

Linear Discriminant Analysis [LDA]



Fisher mapping [fisherm]

PCA is one of the most widely used feature extraction methods

LDA is its supervised cousin

Similar ideas are at the basis of many “novel” methods

PCA



› Principal component analysis [PCA, 1901] :
find directions in data which...

Retain as much [total] variance as possible

Make projected data uncorrelated

Minimize squared reconstruction error

PCA

- › Let a be a projection vector that reduces to 1D
- › Let's say our data has covariance matrix C
 - What value does $a^T C a$ equal to?
 - So, what should we maximize?

PCA

› Let a be a projection vector that reduces to 1D

› Let's say our data has covariance matrix C

What value does $a^T C a$ equal to?


So, what should we maximize?

› Seems we need an assumption...

Assume the constraint $\|a\|^2 = 1$

Then solve, for instance, with Lagrangian : $a^T C a - \lambda(\|a\|^2 - 1)$


PCA Example



› E.g. NIST digits : 2000 samples, $p = 256$

PCA Example

- › For image data, principal components might[!] also be interpretable...
- › Here : largest occurring variations between digits



Remarks on PCA

- › Principal component analysis :

 - Global and linear


 - May need considerable amount of data to estimate covariance $[C, \Sigma, S_T, \dots]$ well

- › Danger :

 - Criterion is not necessarily related to the goal

 - E.g. might discard important directions

 - [Then again, most classifier also do not optimize error rate directly...]



Supervised Linear Feature Extraction

- › If desired output is given,
supervised criteria can be used


One illustration only : Linear Discriminant Analysis
[LDA, or in PRTools terms fisherm]

Intermezzo : Scatter Matrices


- › $m, S_T = \Sigma$: mean and covariance of all samples
- › m_i, Σ_i : mean and covariance of class i
- › Total scatter : Σ equals sum of within and between
- › Within-scatter : $S_w = \sum_{i=1}^C \frac{n_i}{n} \Sigma_i$
- › Between-scatter : $S_B = \sum_{i=1}^C \frac{n_i}{n} (m_i - m)(m_i - m)^T$

Intermezzo : Scatter Matrices

- › S_T = total scatter, “overall width”
- › S_W = “average class width”; the smaller, the better
- › S_B = “average distance between class means”; the larger, the better



LDA [or Fisher mapping]



› Reduction to 1D for two classes

Find projection vector a such that classes are maximally separated

Choose a to maximize Fisher criterion :

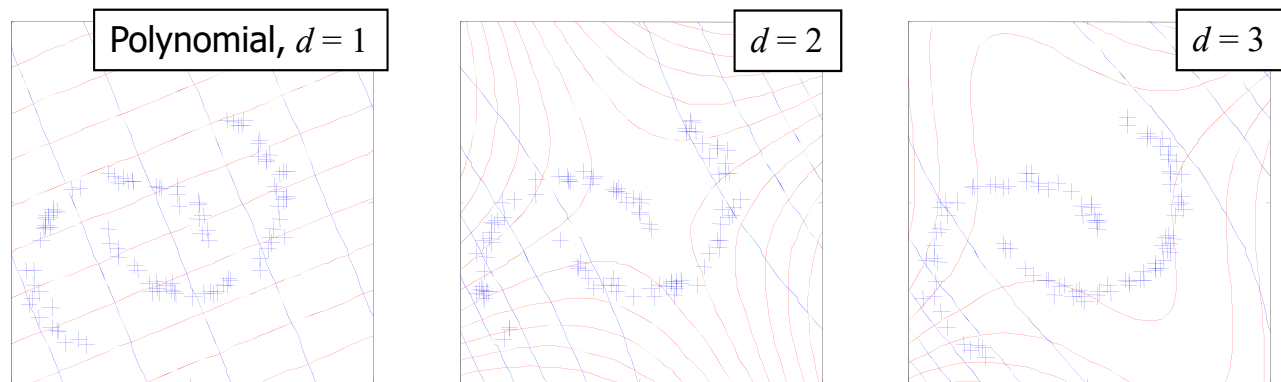
$$J_F(a) = \frac{a^T S_B a}{a^T S_W a}$$

› Solution: Eigenanalysis of $S_W^{-1} S_B$

Nonlinear Feature Extraction

- › Large collection of possible mappings...
- › Today : only one unsupervised method
What is *the* way to make linear stuff nonlinear?
- › Tomorrow : [the gist of] so-called auto-encoders

KPCA



- › Kernelize PCA by relating eigenvectors of $X^T X$ and XX^T [assuming centralized feature vectors]

Summary

- › Feature extraction, like selection :
 - Useful for visualization
 - Necessary because of curse of dimensionality
- › Feature extraction :
 - Linear vs. nonlinear
 - Supervised vs. unsupervised
- › PCA possibly most important method

On to Feature Selection

Feature Selection

- › The general idea is to pick a set of good features from the original/initial set of features

Feature Selection

- › We need a criterion function

 - How do we measure how good a feature subset is?

 - E.g. error, class overlap, information loss

- › We need a search algorithm

 - How do we go through all possible subsets?

 - E.g. pick best single feature at each time

 - › Maybe more than for feature extraction :
optimality is sacrificed

Feature Selection

- › Which approach to feature selection have we seen?

Criteria

Criteria

- › Actual classification performance : “best” possible criterion, but potentially very expensive
- › Approximate performance predictors : calculate easy-to-use measure that gives indication of real performance

Probabilistic Criteria

- › Probabilistic distance
- › Expresses distance between two distributions
 - E.g. Kullback-Leibler divergences and variations
- › Often needs estimates of class-conditional densities
 - Potentially difficult
 - Potentially expensive

Scatter Matrices Again...

- › $\mathbf{m}, S_T = \Sigma$: mean and covariance of all samples
- › \mathbf{m}_i, Σ_i : mean and covariance of class i


› Total scatter : Σ equals sum of within and between

› Within-scatter : $S_w = \sum_{i=1}^C \frac{n_i}{n} \Sigma_i$

› Between-scatter : $S_B = \sum_{i=1}^C \frac{n_i}{n} (\mathbf{m}_i - \mathbf{m})(\mathbf{m}_i - \mathbf{m})^T$

Scatter Matrices Again...

- › S_T = total scatter, “overall width”
- › S_W = “average class width”; the smaller, the better
- › S_B = “average distance between class means”; the larger, the better



Heuristic Scatter-based Criteria

› Example scatter-based performance indicators

$$J_1 = \text{trace} (\mathbf{S}_W + \mathbf{S}_B) = \text{trace} (\mathbf{\Sigma}) = \text{trace}(\mathbf{S}_T)$$

$$J_2 = \text{trace} (\mathbf{S}_B / \mathbf{S}_w)$$

$$J_3 = \det (\mathbf{\Sigma}) / \det (\mathbf{S}_w)$$

$$J_4 = \text{trace} (\mathbf{S}_W) / \text{trace} (\mathbf{S}_B)$$


Yet Another Criterion

- › Mahalanobis distance

$$D_M = (m_1 - m_2)^T C^{-1} (m_1 - m_2)$$

Assumes Gaussian distributions with equal covariance matrix C

In which case, some of the probabilistic distances reduce to this



- › Multi-class, e.g. take sum or minimum

Of course, general solution to extend two-class criteria

- › 1D case : Fisher criterion $J_F = \frac{(m_1 - m_2)^2}{\sigma_1^2 + \sigma_2^2}$

Sub-optimality of Criteria

$$D_M = (m_1 - m_2)^T C^{-1} (m_1 - m_2)$$

- › Give a 2D problem in which Euclidean distance [this assumes $C = I$] picks up the wrong 1D feature...
- › Give a 2D problem in which Mahalanobis distance picks up the wrong 1D feature...

Now, on to Search Algorithms

Now, on to Search Algorithms

- › Feature selection : Select a subset of d out of p measurements which optimizes chosen criterion
- › Simplest solution : look at all possible subsets
Any problems there?

Now, on to Search Algorithms

› Feature selection : Select a subset of d out of p measurements which optimizes chosen criterion

› Simplest solution : look at all possible subsets
Any problems there?

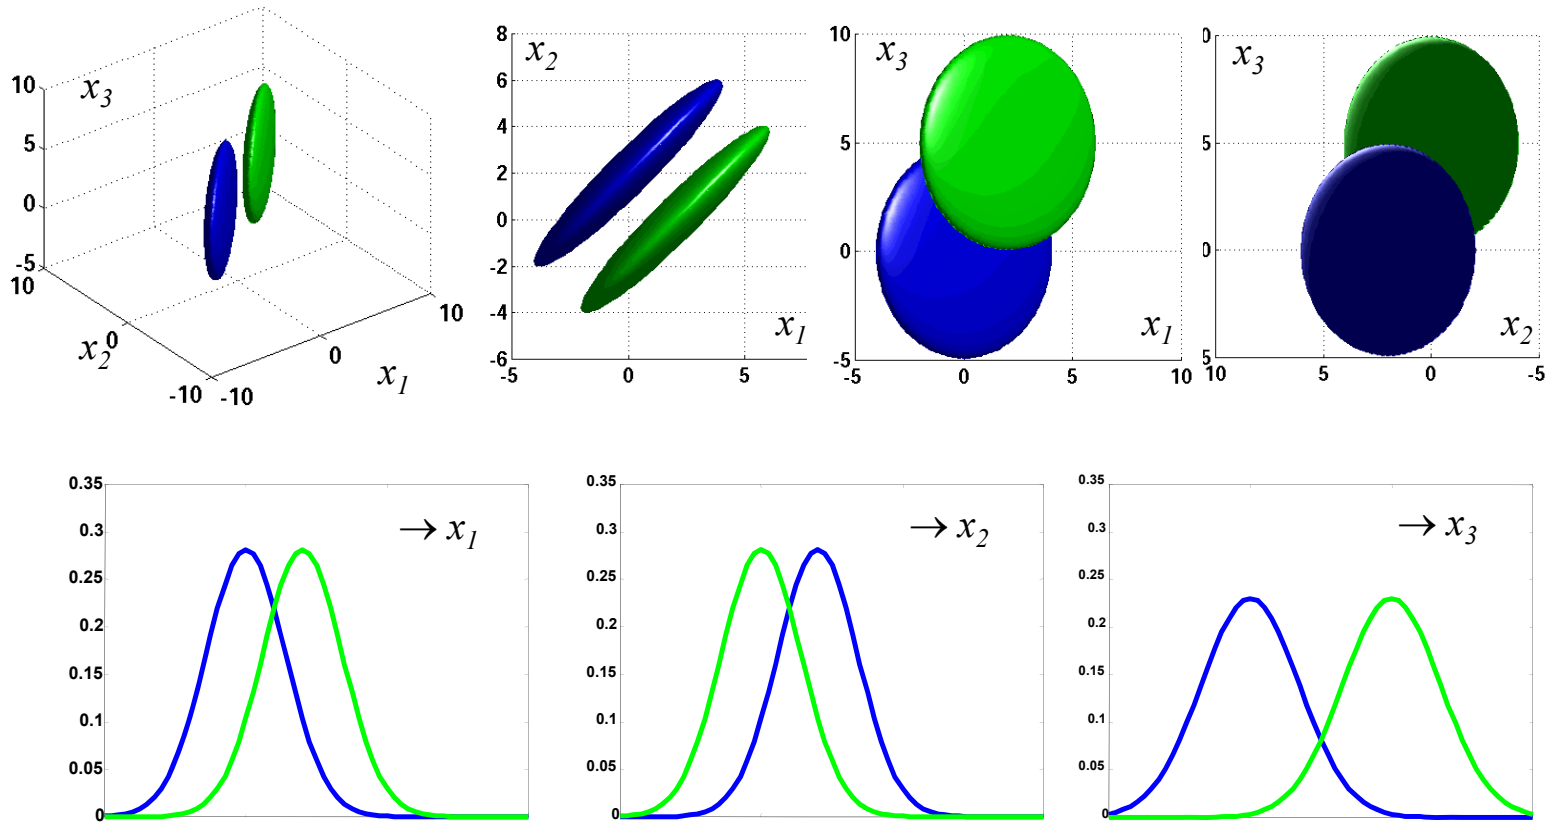
› There are $\binom{p}{d} = \frac{p!}{d!(p-d)!}$ subsets

So, like for the criteria, we settle for approximations...

Search Algorithms

- › Sub-optimal algorithms : select one feature [or a few features] at a time
- › Simplest variation : best individual d
But these are not necessarily the best d !

d Best or Best d ?



More Sub-Optimal Strategies

› Forward selection

Start with empty feature set



One at a time, keep adding feature that gives best performance considering entire chosen feature set

More Sub-Optimal Strategies

› Backward selection

Same as forward selection
...but then the other way 'round

E.g.



- › Select $d = 2$ out of $p = 5$ features

More Sub-Optimal Strategies

› Plus- l -take-away- r



Start with empty set [if $l > r$] or entire set [if $l < r$]

Keep adding best l and removing worst r
[...or vice versa]


Benefit over previous strategies?

Branch & Bound

- › Branch & bound
- › Optimal when criterion is monotonic in number of features d




Floating Search...




- › Backtrack if that improves criterion for given # features

Floating Search...



When Should One Stop?

- › Due to estimation problems [e.g. covariance matrix], criterion may have an optimum
- › Or we could specify desired number of measurements, say, based on data set size
- › Or? Use error rate?



Discussion / Conclusion

- › Some unexpected behaviors in higher dimensions
- › Considered curse of dimensionality
 - Way to counter it and improve performance : lower dimensionality
- › Linear approaches to dimensionality reduction
 - Feature selection and feature extraction
 - Feature extraction can be nonlinear...
- › Approaches are approximative / suboptimal
 - But that holds for many classifiers to start with...

> QQ?