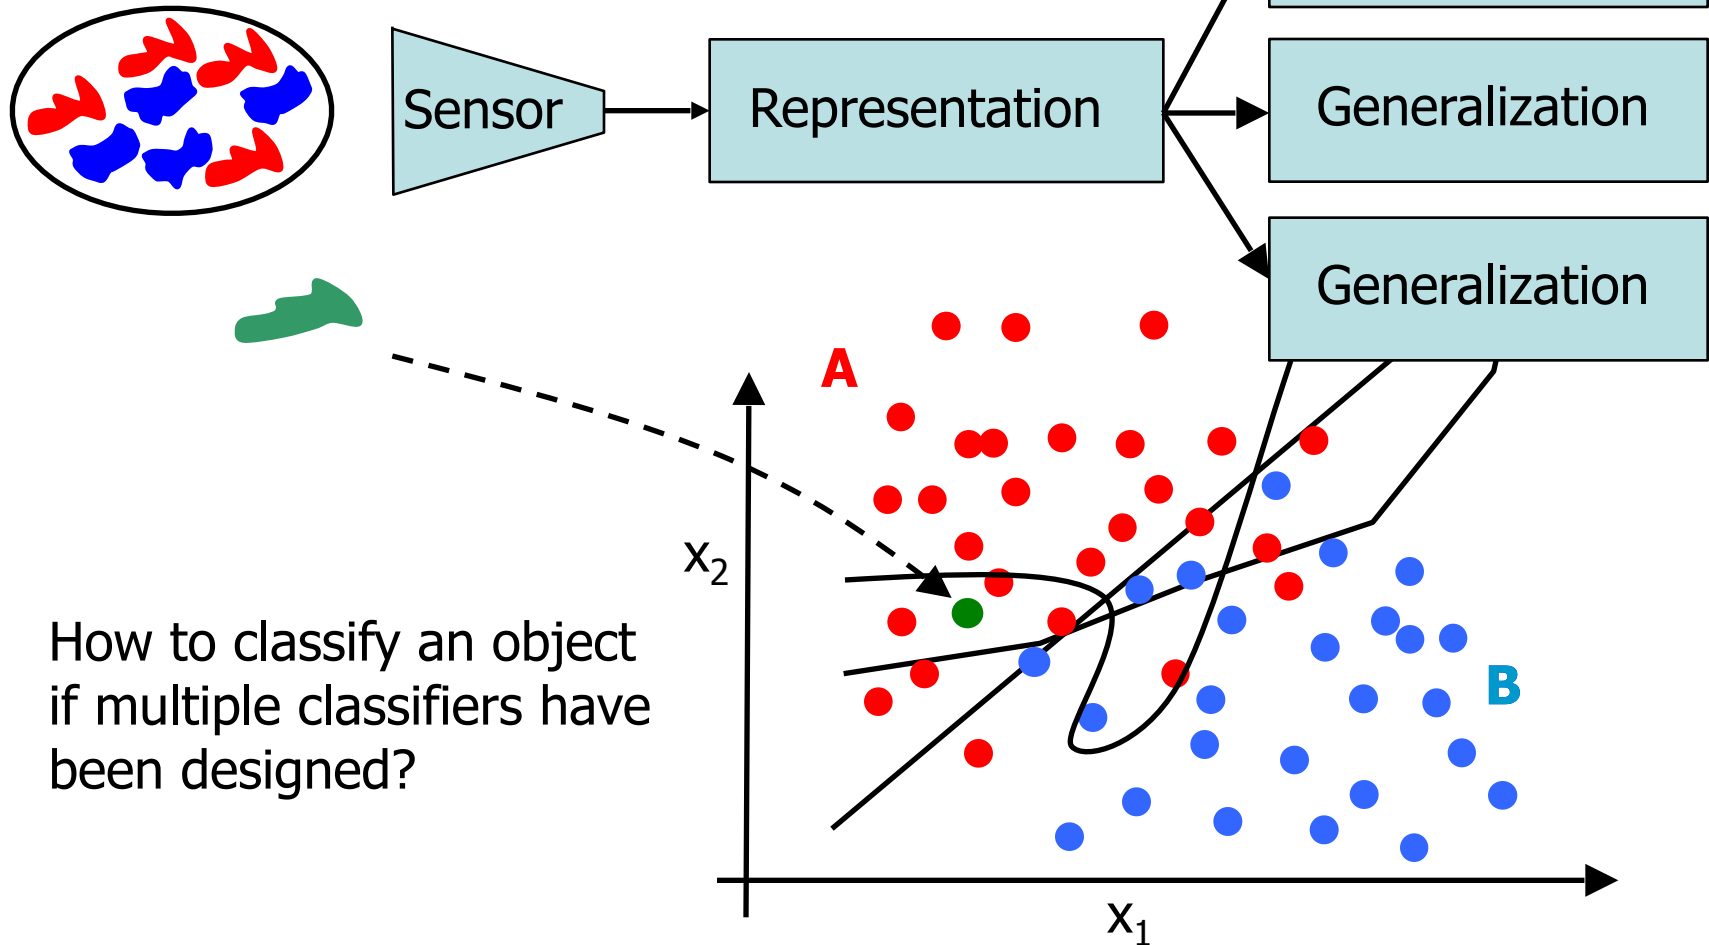


Introduction to Combining Classifiers

D.M.J. Tax



Delft University of Technology

Multiple Classifiers


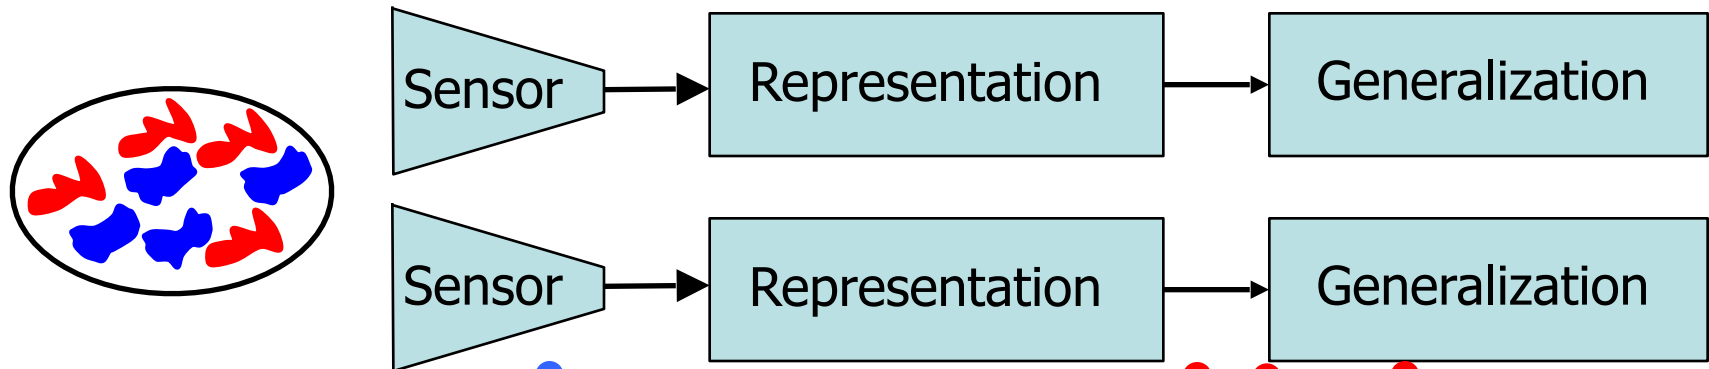


How to classify an object if multiple classifiers have been designed?

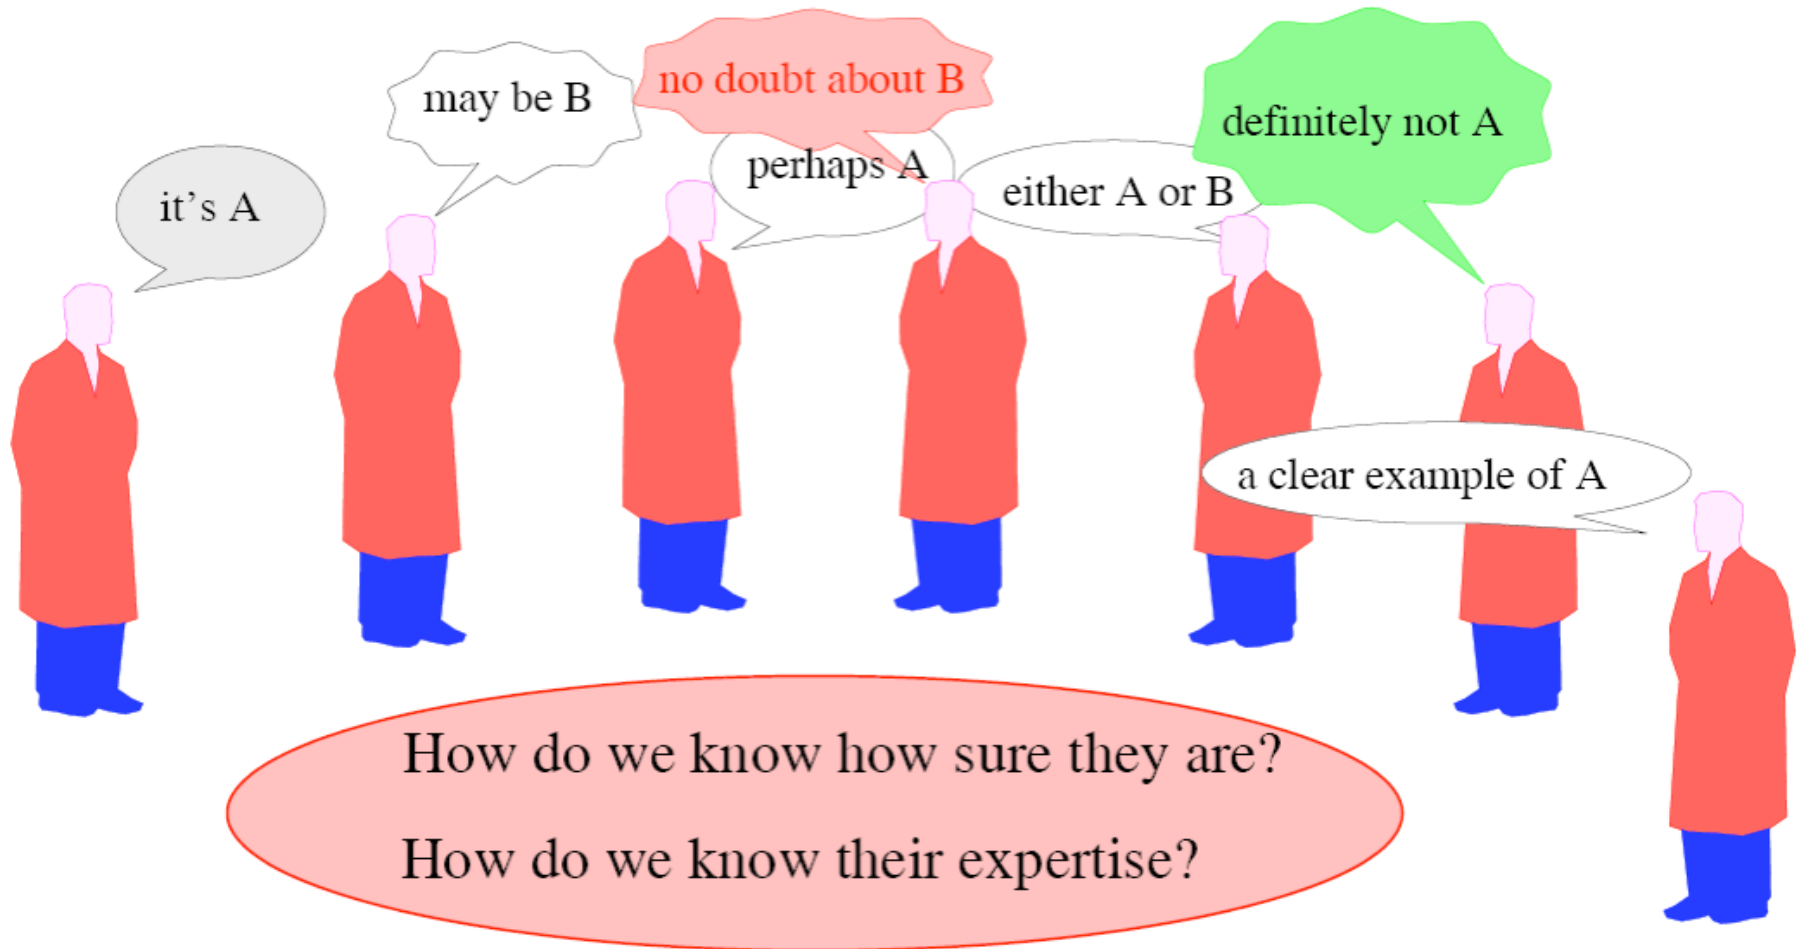
Multiple Representations



Multiple Sensors



Multiple Experts



The Combiner and the Base Classifiers

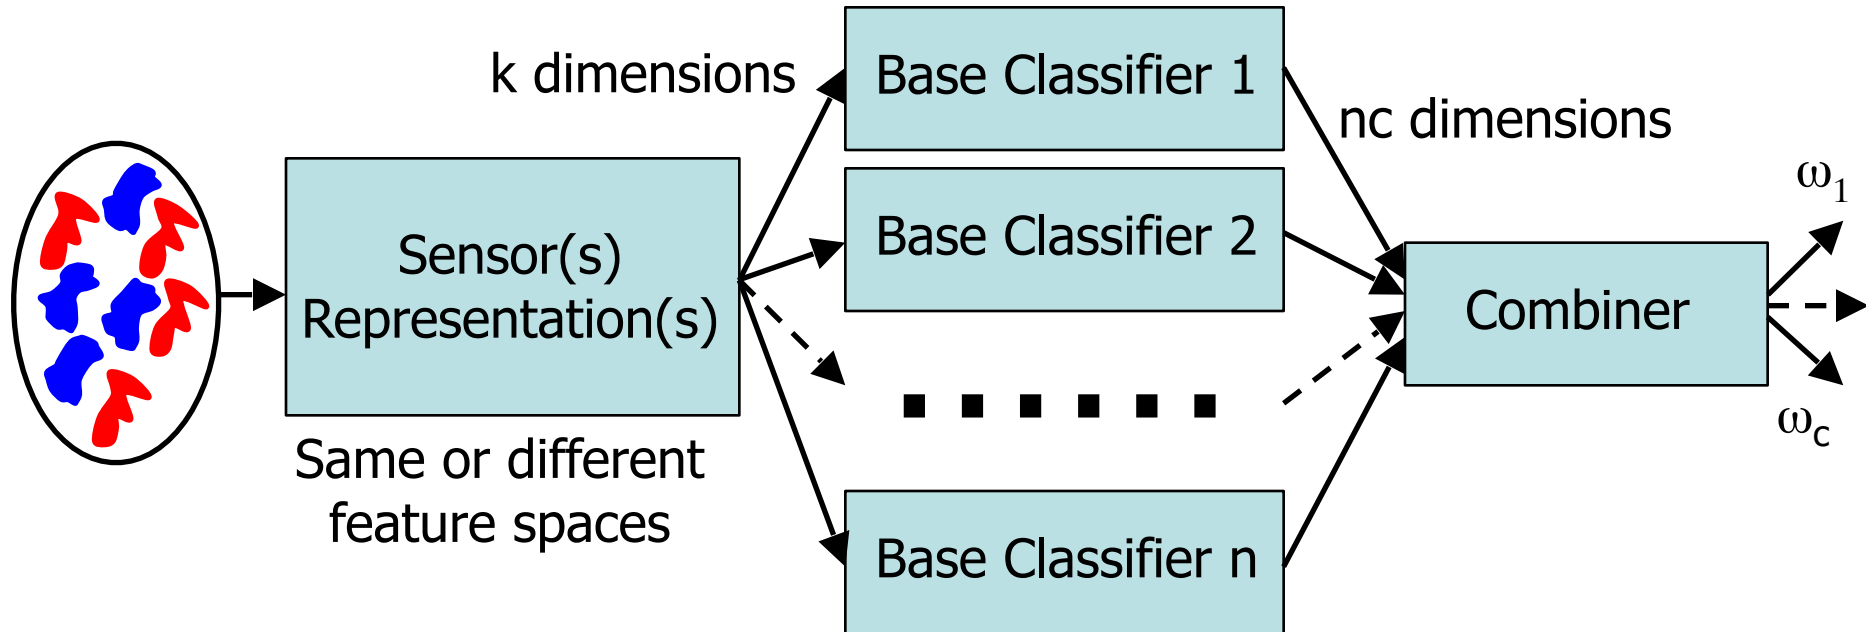
The basic questions:

- How to reach a committee decision?
→ How to design a combining classifier?
- How to constitute a committee?
→ How to generate the base classifiers?

Combining Classifiers

Part I The Combiner

Combining Classifier Architecture



Combining scheme

- Objects: $\{o_i\}$
- Features: $x = F(o)$
User defined representation
- Base classifiers: $y = S(x \mid \text{par_base})$
 par_base : parameters optimized by training set
- Combining classifier $z = C(y \mid \text{par_comb})$
 par_comb : parameters optimized by training set
(sometimes fixed, untrained combiners are used)

Combiners

- Fixed rules based on crisp labels or confidences (estimated posterior probabilities).
- Special trained rules based on 'classifier confidences'.
- General trained rules interpreting base-classifier outputs as features.


Fixed Combining Rules

$$C(y) = \omega_i \text{ if } \operatorname{argmax}_i (\operatorname{comb_rule}_{ji} (y_{ij} = S_{ij}(x_j))) = i$$


An object is assigned to class ω_i if the combination of the outcomes y_{ij} for class ω_i over all classifiers $y_j = S_j(x)$ is maximum.

- | | |
|-----------------------|----------------------|
| ▪ Voting, over labels | $y_{ij} \in \{0,1\}$ |
| ▪ Product, minimum | $y_{ij} \in [0,1]$ |
| ▪ Sum, median | $y_{ij} \in [0,1]$ |
| ▪ Maximum | $y_{ij} \in [0,1]$ |


Voting; Majority Rule, Veto Rule



Confidences, Posterior Probabilities



Posterior Probabilities




Classifier Conditional Posterior Probabilities

$$y_i = S_i(x) = p(\omega_i | x) = \frac{p(\omega_i)f_i(x)}{\sum_i p(\omega_i)f_i(x)}$$


$$y_i \Leftarrow p(\omega_i | S_i(x)) = \frac{p(\omega_i)f_i(y)}{\sum_i p(\omega_i)f_i(y)}$$

Optimal Scaling for Classifier Conditional Posterior Probabilities



Fit a logistic function or sigmoid to the data such that $\prod_{x \in \text{Trainingset}} y(S(x))$ is maximum conditional to $y = 0.5$ for $S(x) = 0$.

Combining Different Representations



Base classifier j posterior probabilities for class A : $y_{Aj} = \text{Prob}_j(A|x_j)$

Product Rule: $y_A = \prod \text{Prob}_j(A|x_j)$, $y_B = \prod \text{Prob}_j(B|x_j)$,

Useful for 'independent' feature spaces (logical 'AND', experts should agree)

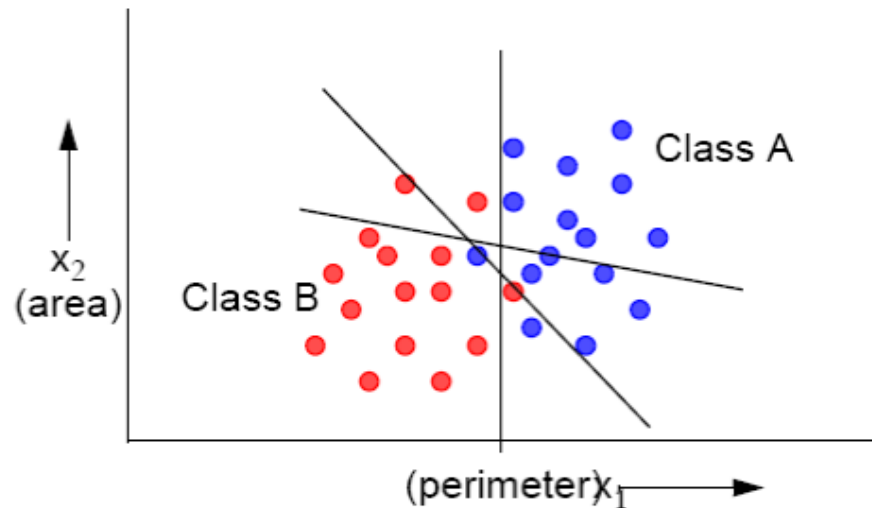
Minimum Rule: $y_A = \text{Min}\{\text{Prob}_j(A|x_j)\}$, $y_B = \text{Min}\{\text{Prob}_j(B|x_j)\}$

Assign according to 'least objecting expert'

prod

min

Combining Different Classifier Estimates



Base classifier j posterior probabilities for class A : $y_{Aj} = \text{Prob}_j(A|\mathbf{x})$

Sum (Mean) Rule: $y_A = \sum \text{Prob}_j(A|\mathbf{x})$, $y_B = \sum \text{Prob}_j(B|\mathbf{x})$,

Useful for improved estimates of posterior probabilities

Also: **Median** and **Majority** Voting

Improvement by averaging out mistakes of experts

mean

median

The Product and the Minimum Rule

Base classifier j posterior probabilities for class A : $y_{Aj} = \text{Prob}_j(A|x_j)$

Product Rule: $y_A = \prod \text{Prob}_j(A|x_j)$, $y_B = \prod \text{Prob}_j(B|x_j)$,

Useful for 'independent' feature spaces, see *Kittler, IEEE-PAMI-20(3), 1998*

Minimum Rule: $y_A = \text{Min}\{\text{Prob}_j(A|x_j)\}$, $y_B = \text{Min}\{\text{Prob}_j(B|x_j)\}$

Assign according to 'least objecting classifier'

objects	Classifier 1		Classifier 2		Product		Minimum	
	Class A	Class B	Class A	Class B	Class A	Class B	Class A	Class B
1	0.4	0.6	0.2	0.8	0.08	0.48	0.2	0.6
2	0.1	0.9	0.7	0.3	0.07	0.27	0.1	0.3
3	0.3	0.7	0.4	0.6	0.12	0.42	0.3	0.6
4	0.5	0.5	0.2	0.8	0.10	0.40	0.2	0.5
5	0.0	1	0.9	0.1	0.00	0.10	0.0	0.1
6	0.8	0.2	0.2	0.8	0.16	0.16	0.2	0.2

Fixed Combining Rules Overview

Product, Minimum

Independent feature spaces

Different areas of expertise

Error free posterior probability estimates

prodc

minc

Ever optimal?

Sum (Mean), Median, Majority Vote

Equal posterior-estimation distributions in same feature space

Differently trained classifiers, but drawn from the same distribution

Bad if some classifiers (experts) are very good or very bad

meanc

majorc

Maximum

Trust the most confident classifier / expert

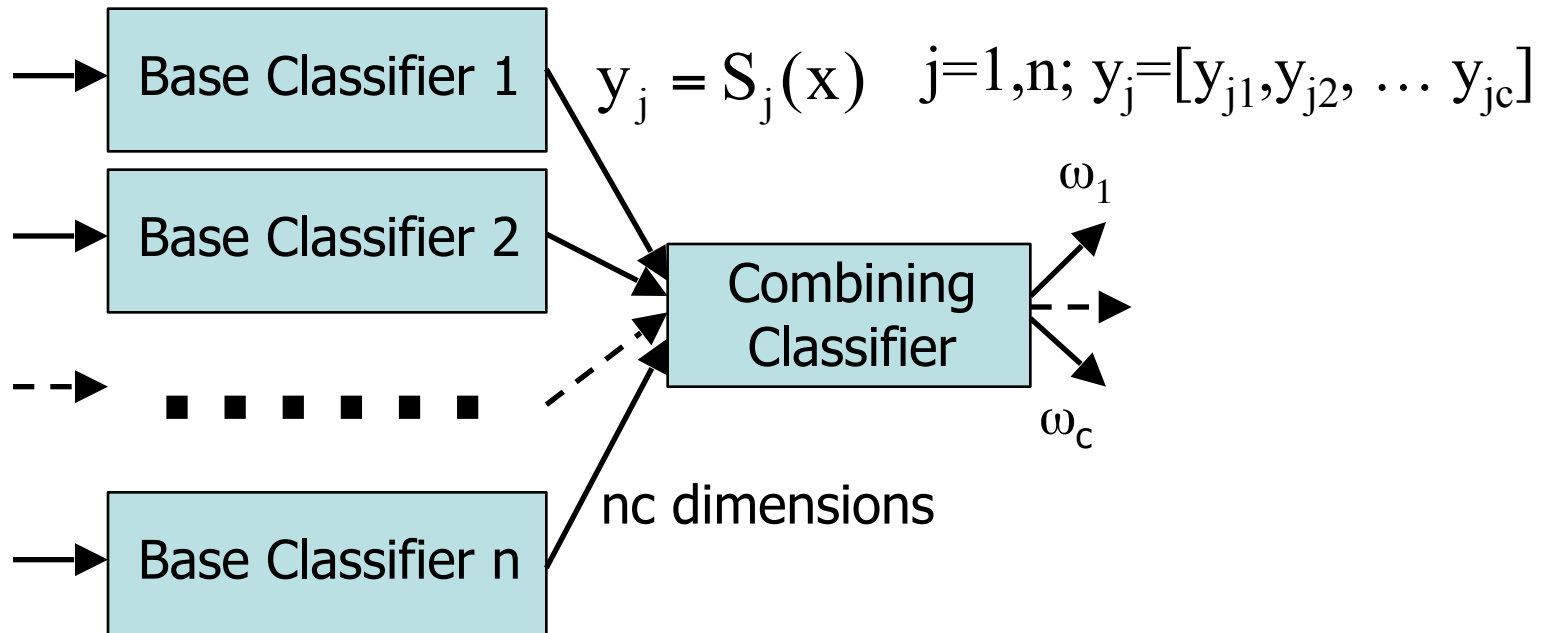
Bad if some classifiers (experts) are badly trained

maxc

Fixed Combining Rules are Sub-optimal

- Base classifiers are never really independent (product)
- Base classifiers are never really equally imperfectly trained (sum, median, majority)
- Sensitivity to over-confident base classifiers (product, min, max)
- Fixed combining rules are never optimal
- Larger training sets do not really improve this (except max?)

Trained Combiners



General rules neglect the classification-confidence characteristic of the base classifier outputs, as they are treated as general feature values.

Trained Combiners (2)

Special Trained Combiners

- DT: Decision Templates (\sim Nearest Mean)
- (BKS: Behavioral Knowledge Space)
- (DCS: Dynamic Classifier Selection)
- (ECOC: Error Correcting Output Coding)
- NN: Neural Networks

General Classifiers


- Nearest Mean
- Fisher
- Decision Trees
-

Decision Templates

- Determine, using a training set, the average outcomes of the base classifiers per class (decision templates, i.e. class means in the base-classifier outcome space).
- Assign new objects to the class of the nearest decision template in the base-classifier outcome space.


(Issue: is it good to have posterior probabilities for the base-classifier outputs, or yield inverse sigmoids a better scaling?)

Behavioral Knowledge Space



Determine on the basis of a training set for every cell in the original feature space the preferred class and assign new objects accordingly.

Dynamic Classifier Selection




- Select the classifier that classifies most of the k nearest neighbors of a test object correctly.
- Use this classifier to classify the test object.

ECOC: Error Correcting Output Coding


- ECOC is a system to use a small set of binary (i.e. two-class) classifiers for a large set of c classes.
- n classifiers can distinguish at most $c=2^n$ classes.
- If $n > 2\log(c)$ the system of classifiers is more robust
- ECOC studies mainly discuss the coding scheme, not the way base classifiers are trained. Combining is done by using the crisp, $\{0,1\}$ -labels.

Example

Combining 10 Bootstrapped Nearest Mean Classifiers




Undertrained, Well Trained, Overtrained



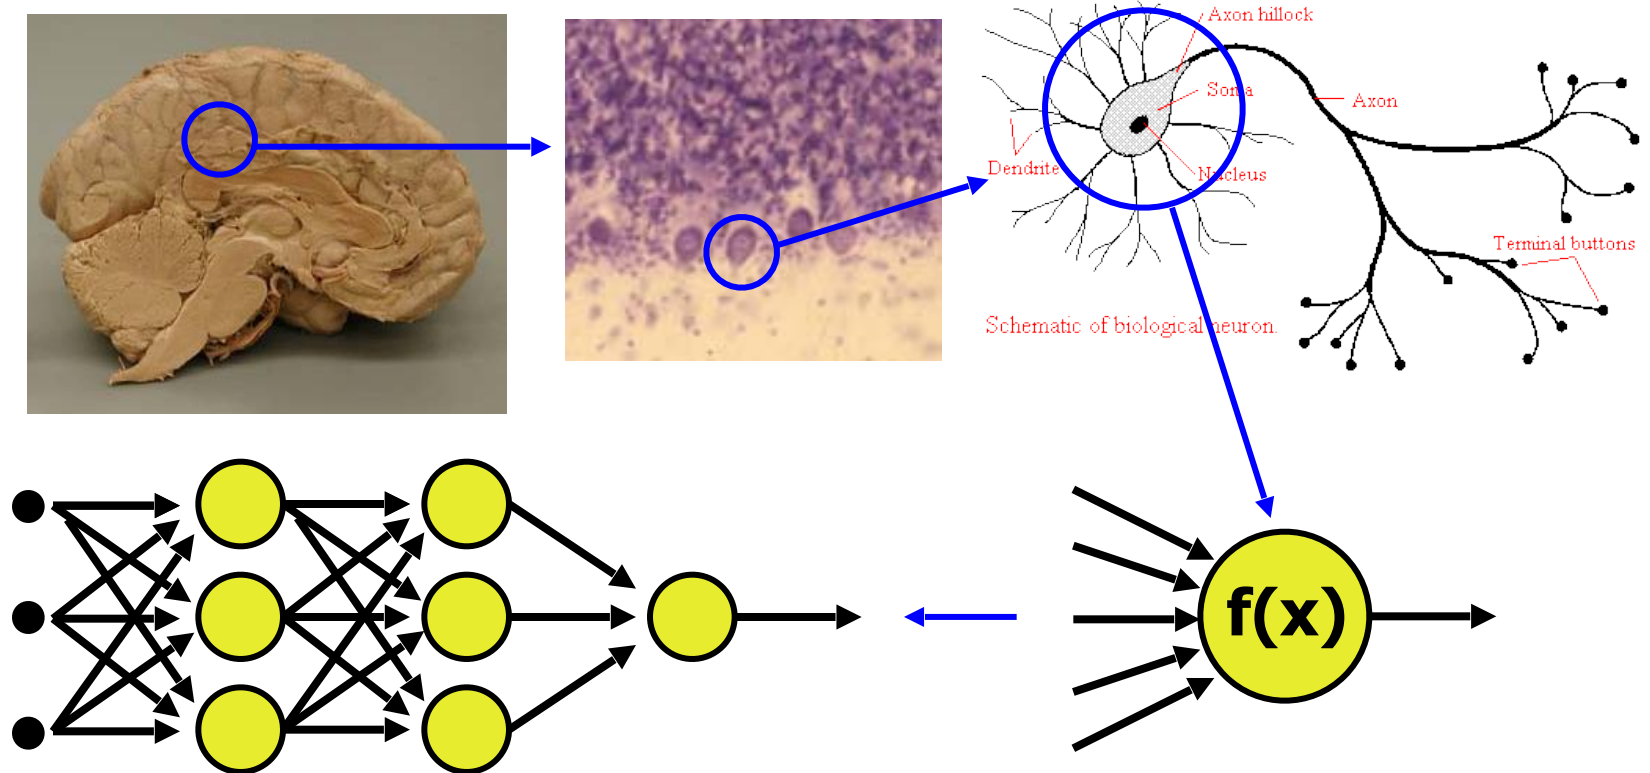
Artificial neural networks (1)

- Large, densely interconnected networks of simple processing units



Artificial neural networks (2)

- Some (not all!) networks originally inspired by the brain




Artificial neural networks (3)

- Research started in the 1950s
- Took off after 1986 – big hype for about 10-15 years
 - brought together psychologists, neurologists, philosophers, machine learners, statisticians...
 - helped thinking about, among others, pattern recognition
 - resulted in a lot of grant money
- Now: useful tool in many areas, with many pitfalls



History

- 1943 : McCulloch and Pitts: model of neuron
- **1958** : Rosenblatt: perceptron
- 1960s : Rosenblatt, Nilsson work on perceptrons
- 1968 : Minsky and Papert point out limitations: perceptrons are linear
- 1982 : Hopfield network (associative memory), Kohonen's self-organising map (clustering), Fukushima's Neocognitron (vision)
- **1986** : Rumelhart, Hinton and Williams: training of nonlinear networks
- 1980s, 1990s: various theoretical developments
- 2010s : More data, faster hardware (GPUs)


McCulloch-Pitts model



Schematic of biological neuron.



McCulloch-Pitts model (2)



weights inputs

$$\text{output } o_i = \phi \left(\sum_j w_{ij} x_j - b_i \right)$$

threshold or bias

with $\phi(a) = \begin{cases} 1 & a \geq 0 \\ 0 & a < 0 \end{cases}$


or $\phi(a) = \frac{1}{1 + \exp(-a)}$



transfer function
or
activation function

Perceptron

- Networks of McCulloch-Pitts models can perform universal computation, given the right weights w : it can do anything a binary computer can do
- ...but how can we find the right weights w ?
- Rosenblatt (1956): possible for single layer networks, perceptrons




Perceptron (2)

- Perceptron is a trainable two-class linear discriminant
- Training algorithm can be proven to converge to correct solution for separable classes
- Possible to extend to multiple classes
- When classes are not linearly separable:
 - indefinite training, weights will blow up
 - solution: decrease ρ during training

Perceptron (3)

- Minsky & Papert (1969): perceptrons are limited




The XOR problem
cannot be
solved by a linear
discriminant
such as the perceptron

- When classes are nonlinearly separable:
 - nonlinear transfer functions
 - multilayer perceptron – but how to find weights...?
 - Rumelhart et al. (1986): use the chain rule

Multilayer perceptron


- Stacked perceptrons: feedforward networks
- Each unit has a nonlinear transfer function,

e.g. $\phi(a) = \frac{1}{1 + \exp(-a)}$ sigmoid or logistic function



Backpropagation training


- To explain training: simple network
- One hidden layer, linear output units



- $$a_j^o = \sum_{i=1}^{n_{o-1}} w_{ji}^o z_i^{o-1}, \quad z_j^{o-1} = \phi(a_j^{o-1})$$

Other training algorithms

- Backpropagation training is simple gradient descent, but implemented in a useful way: all updates can be calculated locally (in parallel)
- Other view: simply optimise MSE E w.r.t. weight vector w



using any optimisation routine, e.g.

- second order (Newton, pseudo-Newton)
- conjugate gradient descent
- Broyden-Fletcher-Goldfarb-Shanno (BFGS)
- Levenberg-Marquardt (LM, in PRTools)

Multilayer perceptrons (2)


- Choices:
 - targets (0/1, 0.1/0.9, 0.2/0.8) t
 - **number of hidden layers**
 - **number of units per hidden layer** n_i
 - transfer functions $\phi(a)$
 - initialisation $w^{(0)}$
 - training algorithm
 - parameters (learning rate ρ etc.)
 - convergence decision E_{thr} or test set selection
 - ...
- All of these influence results!

Multilayer perceptrons (3)


- Number of weights = number of parameters = $\sum_{l=1}^{o-1} (n_l + 1)n_{l+1}$
e.g. for $p = 10$, $C = 2$, 2 20-unit hidden layers:
 $(10 + 1) \cdot 20 + (20 + 1) \cdot 20 + (20 + 1) \cdot 2 = 682$ parameters
- Danger of overtraining!
- Prevention:
 - use small networks
 - regularise: minimise $E(\mathbf{w}) + \lambda \|\mathbf{w}\|$
 - small w 's: low complexity, training slowly increases w 's;
so when stopping in time: automatic regularisation!

Multilayer perceptrons (4)

- Examples: 1 hidden layer of 3 units, 2 initialisations



2 hidden layers of 5 units each, 2 initialisations



Combining Classifiers

Part II

Generation of Base Classifiers

Generation of base classifiers

- Random Subspace Approach
- Bagging
- Boosting

Random Subspace Approach

- Select a dimensionality $k' \ll k$ that fits well with the training size
 - Select at random n subsets of k' features
 - Train n classifiers
 - Combine
-
- When better than feature selection?
 - When better than feature extraction?


Bagging (Bootstrap and Aggregate)

- Select a training set size $m' < m$
 - Select at random n subsets of m' training objects (original: bootstrap)
 - Train a classifier (original: decision tree)
 - Combine (original: vote)
-
- Bagging decision trees combined with random subspace approach gives: Random Forest classifier


Boosting

- Initialize all objects with an equal weight
- Select a training set size $m' < m$ according to the object weights
- Train a weak classifier
- Increase the weights of the erroneously classified objects
- Repeat as long as needed
- Combine


Boosting: Emphasize Difficult Objects




Dataset




Step 1




Step 2




Step 3




Step 4




Step 5




Step 6




Step 7




Step 8



Step 9



Step 10




Final

Adaboost Algorithm

1. **Sample** the training set according to a set of object weights (initially equal)
2. Use it for **training a** simple (weak) **classifier** w_i
3. **Classify** the entire data set, using the weights, **error** ϵ_i
Store classifier weight $\alpha_i = 0.5 \log((1-\epsilon_i)/\epsilon_i)$
4. **Multiply weights** of erroneously classified objects with $\exp(\alpha_i)$
Multiply weights of correctly classified objects with $\exp(-\alpha_i)$
5. **Goto 1** as long as needed
6. **Final classifier**: weighted voting, weights α_i


Assumes decision stump as weak classifier, and minimises exponential loss (See course CS4230 Machine Learning 2)

Adaboost - 2D Example




100 dec. stumps, wvote

```
w=adaboostc(a,stumpc,100)
```



10 dec. stumps, Fisher

```
w=adaboostc(a,stumpc,10,fisher)
```



10 Fisher, Fisher

```
w=adaboostc(a,fisher,10,fisher)
```

Discussion on Base Classifiers

- Are to be combined
- Simple, not overtrained, especially not for trained combiners
- Many: fast training, fast execution
- Soft outputs might be helpful
- Traditional: decision trees, decision stumps, linear, quadratic
- Weak classifiers: simple, should do something,
 - not sufficient for the problem,
 - large bias, large variance

Discussion on Combining

- Base classifiers are trained on systematically different training sets.
- They may be different in importance → weighting is appropriate.

$$S(x) = \text{sign}\left(\sum_i \alpha_i w_i(x)\right)$$

- Weights should follow from the performance → training
- Weights are not optimized for the ensemble of classifiers (avoid overtraining)
- Weights operate on the crisp outputs of $w_i(x)$: weighted voting
- Alternative:
Fisher: soft outputs, optimize weights over ensemble of base classifiers
Large set of base classifiers may overtrain